We sincerely thank all the reviewers for their valuable feedback and insightful comments. We are glad that the reviewers appreciated the motivation of our work, the VISDA dataset, and the proposed evaluation framework for studying situational effects under fixed value profiles. We have carefully considered all the suggestions given by the reviewers.

The reviewers raised a few common concerns regarding model coverage, the number of base scenarios, and the scalability of the human study. We have addressed these concerns by providing additional experimental evidence and clarifications. Below, we respond to each reviewer's comments individually.

Response to Reviewer 5e7B

We sincerely thank the reviewer for the detailed and constructive feedback.

Comment 1: "*Limited model coverage*"

We thank the reviewer for the above observation. Our initial evaluation focused on open-weight models to ensure reproducibility and enable the community to replicate all experiments without relying on proprietary APIs. However, after the ARR submission, we have additionally evaluated **Claude Haiku 4.5**, a commercial frontier model. We observed an overall **decision-flip rate of 10.86%**, which is consistent with the trends reported for Gemma (8.92%), Qwen (6.58%), and LLaMA (2.03%). This additional experiment suggests that the observed situational sensitivity is not limited to the originally evaluated open-weight models. We will include these new experimental results and discussion in the revised manuscript.

Comment 2: "*Only 10 base scenarios*"

We agree that increasing the number and diversity of base scenarios would further strengthen the observation on VISDA . Our current design prioritizes **controlled evaluation** by systematically varying value profiles and situational modifiers while keeping the underlying scenario fixed, allowing us to isolate the effect of contextual changes.

While VISDA contains **10 base scenarios**, these are systematically expanded across **8 situational modifier axes** and **95 Schwartz value profiles**, resulting in (10x8x95) **7,600 modifier evaluations per model**. Expanding the base scenarios by simply doubling the number of base scenarios would proportionally increase the evaluation space to **15,200 evaluations per model**, while preserving the same evaluation protocol. Since our experiments are conducted across multiple LLMs, this substantially increases both computational cost and evaluation time.

The primary objective of this work was to establish a controlled and reproducible benchmark rather than maximize the number of scenarios. That said, we agree that expanding the set of base scenarios would further strengthen the robustness of our observations. We will conduct the experiment with 20 base points and will include the results in the final version of the paper upon acceptance.

Comment 3: “*Selection of the 95 value profiles*”

Appendix A and section 3.1 in the paper describe the value profile selection process in detail and also highlight additional information. We will release the exact value profiles used in the experiment along with the code and the dataset as mentioned in the paper (cf. page 2 footnote).

Comment 4: “*Modifier validation*”

We appreciate this important suggestion. Although all generated scenarios and modifiers were manually reviewed for plausibility and consistency before inclusion in VISDA, we agree that this validation process was not sufficiently described in the manuscript.

To address this concern, we conducted a dedicated human validation study in which independent evaluators assessed the generated modifiers with respect to (i) semantic preservation of the original scenario and (ii) correctness of the intended situational modifier axis. The validation results showed strong agreement, indicating that the generated modifiers consistently preserved the underlying decision problem while accurately representing the intended contextual pressure. These observations are consistent with the assumptions underlying our evaluation methodology.

We will include the validation protocol and summary statistics in the revised manuscript to improve transparency. Furthermore, as noted in the footnote on Page 2, we will publicly release the dataset and evaluation code upon acceptance.

We thank the reviewer once again for the thoughtful feedback. We believe these revisions substantially improve the clarity, transparency, and experimental strength of the paper while preserving its central contribution: introducing VISDA as a controlled benchmark for evaluating how situational context influences value-conditioned moral decisions in both humans and LLMs.

Response to Reviewer 19NF

We sincerely thank the reviewer for the thorough assessment of our work and for recognizing the potential value of VISDA as a benchmark for studying situational robustness in moral reasoning. We are encouraged that the reviewer found the proposed dataset and the overall findings interesting. We appreciate the constructive suggestions and address them below.

Comment 1: "*Limited number of base scenarios*"

This concern is also raised by another reviewer. We agree that increasing the number and diversity of base scenarios would further strengthen the observation on VISDA . Our current design prioritizes **controlled evaluation** by systematically varying value profiles and situational modifiers while keeping the underlying scenario fixed, allowing us to isolate the effect of contextual changes.

While VISDA contains **10 base scenarios**, these are systematically expanded across **8 situational modifier axes** and **95 Schwartz value profiles**, resulting in (10x8x95) **7,600 modifier evaluations per model**. Expanding the base scenarios by simply doubling the number of base scenarios would proportionally increase the evaluation space to **15,200 evaluations per model**, while preserving the same evaluation protocol. Since our experiments are conducted across multiple LLMs, this substantially increases both computational cost and evaluation time.

The primary objective of this work was to establish a controlled and reproducible benchmark rather than maximize the number of scenarios. That said, we agree that expanding the set of base scenarios would further strengthen the robustness of our observations. We will conduct the experiment with 20 base points and will include the results in the final version of the paper upon acceptance.

Comment 2: "*Small and non-representative human sample*"

We agree that our participant pool is relatively small and not representative of the broader population. However, the objective of the human study was not to estimate population-level human values, but rather to provide a **pilot comparison** with LLMs using the proposed **decision-flip metric** under controlled situational modifiers. In this regard, the study demonstrates that humans, like LLMs, also exhibit decision changes under contextual pressure and provides a meaningful reference point for comparison.

We agree that increasing both the size and demographic diversity of the participant pool would further strengthen the study. However, recruiting a large and well-balanced population across multiple demographic groups is practically challenging. Although increasing diversity by including participants from medicine, engineering, software, education, and other professions broadens the evaluation, it often results in relatively small sample sizes within each subgroup. Consequently, statistically meaningful subgroup analyses require a substantially larger overall study population. We will revise the manuscript to more clearly position the current human study as a pilot evaluation, explicitly discuss these limitations, and identify a larger and more diverse human study as an important direction for future work.

We thank the reviewer once again for the encouraging feedback and constructive suggestions. We believe these revisions improve the clarity, transparency, and overall strength of the paper while preserving its primary contribution of introducing VISDA as a controlled benchmark for evaluating situational effects in value-conditioned moral decision-making.

Response to Reviewer Xqnr

We sincerely thank the reviewer for the encouraging and thoughtful feedback. We are glad that the reviewer appreciated the motivation of our work, recognized the limitations of existing value-based evaluations, and found VISDA to be a useful and extensible benchmark. We also appreciate the positive comments on the utility-maximization baseline and the modular dataset design.

Comment 1: "*Limited model evaluation*"

This concern is also raised by another reviewer. Our initial evaluation focused on reproducible open-weight models to ensure reproducibility and enable the community to replicate all experiments without relying on proprietary APIs. However, after the ARR submission, we have additionally evaluated **Claude Haiku 4.5**, a commercial frontier model. We observed an overall **decision-flip rate of 10.86%**, which is consistent with the trends reported for Gemma (8.92%), Qwen (6.58%), and LLaMA (2.03%). This additional experiment suggests that the observed situational sensitivity is not limited to the originally evaluated open-weight models. We will include these new experimental results and discussion in the revised manuscript.

Comment 2: "*Human study*"

This concern is also raised by another reviewer. We appreciate the reviewer's observation regarding the size of the human study. As noted in the paper, our intention was to use the human experiment as a **pilot comparison** with LLM behavior using the proposed decision-flip metric, rather than to make population-level claims about human moral decision-making. We agree that a larger and more diverse participant pool would further strengthen the study, and we will make this limitation more explicit in the revised manuscript while identifying larger-scale human evaluation as an important direction for future work.

We thank the reviewer once again for the positive assessment and constructive suggestions. We believe the additional experiments, clearer presentation, and expanded discussion will further strengthen the paper while preserving its primary contribution of introducing VISDA as a controlled benchmark for evaluating situational effects in value-conditioned moral decision-making.

We sincerely thank the reviewers for their time and valuable feedback. We believe that the additional experiments, methodological clarifications, and planned revisions substantially strengthen the paper and address the major concerns raised during the review process. We kindly request the reviewers to reconsider their evaluation and review scores in light of these responses.